



Current best architecture · CIFAR-10

32-convolution residual CNN | 1,239,274 trainable parameters



Inside each residual block



$$y = \text{ReLU}(\text{BN}_2(\text{Conv}_2(\text{ReLU}(\text{BN}_1(\text{Conv}_1(x)))))) + S(x)$$

Both convs: stride 1, padding 1, no bias
S(x): identity; append zeros when channels increase

Validation: 87.20 ± 0.57% · Test: 86.43 ± 0.49% · 3 seeds, 160 epochs, no augmentation