


CIFAR-10 CNN

Current network architecture · 545,098 trainable parameters



Flow: bottom → top. Shapes are channels × height × width; batch dimension omitted.
CrossEntropyLoss consumes the logits directly. Convolutions have stride 1; pooling has stride 2.