




Figure 1: The MLP architecture. Trained with AdamW (lr 2×10^{-3} , weight decay 0.05), one-cycle LR schedule, label smoothing 0.1, batch 128, 50 epochs. 1.87M parameters; 99.41% MNIST test accuracy.